

A Survey on GANs (Cycle GAN, WGAN, Style GAN)

Amirmahdi Ansaripour

Table of Contents

- Cycle GAN
- Wasserstein GAN (WGAN)
- Style GAN
- References

Cycle GAN

- It was first introduced in the “Unpaired Image-to-Image Translation Using Cycle Consistent Adversarial Networks” paper.
- As the name of the paper suggests, it is used for image-to-image translation; for example, between colored and grayscale images, edge-map and photograph, zebras and horses, etc.



Figure 1: Having two source domains (zebras, horses) or (Monet, photo), Cycle GAN can convert images from one domain to the other one.

Cycle GAN

- How does the model work?

It uses two GAN structures, in other words, two generators and two discriminators. The generators are shown as G and F .

Mathematically, they can be written as:

$$G: X \rightarrow Y$$

$$F: Y \rightarrow X$$

The discriminators, D_X and D_Y , try to distinguish between the transformed images and real images. Our objective is to optimize the components such that:

$$\begin{aligned} F(G(X)) &\approx X \\ G(F(Y)) &\approx Y \end{aligned} \tag{1}$$




Figure 2: A scheme of Cycle GAN.

Cycle GAN

- Loss functions:

1. Typical loss functions of Vanilla GAN: The same is for F and D_X .

$$\begin{aligned}\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]\end{aligned}\tag{2}$$

1. Cycle Consistency Loss: Norm-1 of the distance between $F(G(X))$ and X and $G(F(Y))$ and Y . Please refer to Eq. 1.

$$\begin{aligned}\mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].\end{aligned}\tag{3}$$

Cycle GAN

- The final loss function then becomes:

$$\begin{aligned}\mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F),\end{aligned}\tag{4}$$

Cycle GAN

- Cycle loss implementation:

```
# Cycle loss
map_to_B = generator_to_B(real_train_A, training = True)
cycled_A = generator_to_A(map_to_B, training = True)
map_to_A = generator_to_A(real_train_B, training = True)
cycled_B = generator_to_B(map_to_A, training = True)
total_cycle_loss = calc_cycle_loss(real_train_A, cycled_A) + calc_cycle_loss(real_train_B, cycled_B)
```

- Vanilla GAN losses implementation: Note that we apply each discriminator on each category of images (mapped and real), and there are two discriminators.

```
# Discriminator loss
disc_A_real = discriminator_on_A(real_train_A, training = True)
disc_A_generated = discriminator_on_A(map_to_A, training = True)
disc_B_real = discriminator_on_B(real_train_B, training = True)
disc_B_generated = discriminator_on_B(map_to_B, training = True)
disc_A_loss = calc_discriminator_loss(real = disc_A_real, generated = disc_A_generated)
disc_B_loss = calc_discriminator_loss(real = disc_B_real, generated = disc_B_generated)
# Generator loss
gen_A_loss = calc_generator_loss(generated = disc_A_generated)
gen_B_loss = calc_generator_loss(generated = disc_B_generated)
```

Cycle GAN

Some sample outputs (Zebra \leftrightarrow Horse domains).

Interesting point in this pair



Figure 3: Pairs of (source, transformed) images obtained in my implementation.

Table of Contents

- Cycle GAN
- Wasserstein GAN (WGAN)
- Style GAN
- References

Wasserstein GAN (WGAN)

- It was first introduced in the “**Wasserstein GAN**” paper, and then improved in the “**Improved Training of Wasserstein GANs**”.
- It criticizes the KL¹ and JS² divergences, stating that they are not continuous and differentiable in all points. For example:

let $Z \sim U[0, 1]$ and P_0 be the distribution of $(0, z) \in R^2$, a line from point $(0, 0)$ to $(0, z)$.

Also, P_θ is the distribution of $(\theta, z) \in R^2$, a line from point $(\theta, 0)$ to (θ, z) .

$$JS(\mathbb{P}_0, \mathbb{P}_\theta) = \begin{cases} \log 2 & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases} \quad KL(\mathbb{P}_0 \| \mathbb{P}_\theta) = \begin{cases} +\infty & \text{if } \theta \neq 0, \\ 0 & \text{if } \theta = 0, \end{cases} \quad (5)$$

WGAN

- As we can see, both JS and KL divergences are non-continuous, so using them as objectives is not always beneficial.
- The paper proposes Wasserstein-1 distance, also called Earth-Mover's distance, instead of previous divergences.

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|], \quad (6)$$

In the above formula, inf means infimum or the greatest lower bound, and $\Pi(P_r, P_g)$ is the set of all joint distributions $\gamma(x, y)$. The Wasserstein loss between the P_0 and P_θ distributions is therefore:

$$W(\mathbb{P}_0, \mathbb{P}_\theta) = |\theta|. \quad (7)$$




Figure 4: JS divergence (top) and Wasserstein loss (bottom) between P_0 and P_θ in the previous slide.

WGAN

- Intuitively, $W(P_r, P_g)$ states how much “mass” should be transported from x to y in order to transform the distribution P_r (real) into P_g (generated). WGAN tries to minimize this mass or at least find its greatest lower bound.
- However, finding the infimum in Eq. 6 is intractable, and the paper introduces solving a duality problem instead:

$$\max_{\|f\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r}[f(x)] - \mathbb{E}_{x \sim \mathbb{P}_\theta}[f(x)] \quad (8)$$

In Eq. 8, f is 1-Lipschitz function, which is estimated by a neural network:

$$\frac{|f(x) - f(y)|}{|x - y|} \leq 1 \quad (9)$$

WGAN

- The final algorithm for training WGAN
- Two points:
 1. The discriminator in Vanilla GAN is named critic, whose duty is to estimate f with weights w .
 2. In each epoch, first, the critic is trained in n_{critic} inner loops, and then the generator is trained.
 3. Line 7 enforces f to be 1-Lipschitz. The paper discusses that c is a hyper-parameter and in practice should take values bigger than 1. It highlights the disadvantages of setting c to very low or high values.

Algorithm 1 WGAN, our proposed algorithm. All experiments in the paper used the default values $\alpha = 0.00005$, $c = 0.01$, $m = 64$, $n_{critic} = 5$.

Require: : α , the learning rate. c , the clipping parameter. m , the batch size. n_{critic} , the number of iterations of the critic per generator iteration.

Require: : w_0 , initial critic parameters. θ_0 , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{critic}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w [\frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSProp}(\theta, g_\theta)$ 
12: end while
```

Figure 5: Training algorithm of WGAN with simple clipping.

WGAN

- Enforcing f to be 1-Lipschitz through clipping is a very naive approach. In a second paper named “Improved Training of Wasserstein GANs”, a penalty term is introduced when training the critic (Line 7).

Algorithm 1 WGAN with gradient penalty. We use default values of $\lambda = 10$, $n_{\text{critic}} = 5$, $\alpha = 0.0001$, $\beta_1 = 0$, $\beta_2 = 0.9$.

Require: The gradient penalty coefficient λ , the number of critic iterations per generator iteration n_{critic} , the batch size m , Adam hyperparameters α, β_1, β_2 .
Require: initial critic parameters w_0 , initial generator parameters θ_0 .

```
1: while  $\theta$  has not converged do
2:   for  $t = 1, \dots, n_{\text{critic}}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\tilde{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$ 
6:        $\hat{\mathbf{x}} \leftarrow \epsilon \mathbf{x} + (1 - \epsilon) \tilde{\mathbf{x}}$ 
7:        $L^{(i)} \leftarrow D_w(\tilde{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda(\|\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while
```

Figure 6: Training algorithm of WGAN with gradient penalty term method.

WGAN




Figure 7: My output generated samples on CIFAR-10 (left) and MNIST (right) datasets.

Table of Contents

- Cycle GAN
- Wasserstein GAN (WGAN)
- Style GAN
- References

Style GAN

References

- [1] Zhu, J.Y., Park, T., Isola, P. and Efros, A.A., 2017. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision* (pp. 2223-2232).
- [2] Arjovsky, M., Chintala, S. and Bottou, L., 2017, July. Wasserstein generative adversarial networks. In *International conference on machine learning* (pp. 214-223).
- [3] Gulrajani, I., Ahmed, F., Arjovsky, M., Dumoulin, V. and Courville, A.C., 2017. Improved training of wasserstein gans. *Advances in neural information processing systems*, 30.